

COMP6771

Advanced C++ Programming

Week 5.1

Smart Pointers

A recap on pointers

Sy Brand explains pointers in two minutes by feeding their cat!

Object lifetimes

To create safe object lifetimes in C++, we always attach the lifetime of one object to that of something else

- Named objects:
 - A variable in a function is tied to its scope
 - A data member is tied to the lifetime of the class instance
 - An element in a `std::vector` is tied to the lifetime of the vector
- Unnamed objects:
 - A heap object should be tied to the lifetime of whatever object created it
 - Examples of bad programming practice
 - An **owning raw pointer** is tied to nothing
 - A **C-style array** is tied to nothing
- **Strongly recommend** watching the first 44 minutes of Herb Sutter's cppcon talk "**Leak freedom in C++... By Default**"

Object lifetime with references

We need to be very careful when returning references.

The object must always outlive reference.

```
auto okay(int& i) -> int& {  
    return i;  
}  
  
auto okay(int& i) -> int const& {  
    return i;  
}
```

```
auto questionable(int const& x) -> int const& {  
    return i;  
}  
  
auto not_okay(int i) -> int& {  
    return i;  
}  
  
auto not_okay() -> int& {  
    auto i = 0;  
    return i;  
}
```

Creating a safe* pointer type

Don't use the new / delete keyword in your own code
We are showing for demonstration purposes

```
1 // myintpointer.h
2
3 class MyIntPtr {
4     public:
5         // This is the constructor
6         MyIntPtr(int* value);
7
8         // This is the destructor
9         ~MyIntPtr();
10
11     int* GetValue();
12
13     private:
14         int* value_;
15 };
```

```
1 // myintpointer.cpp
2 #include "myintpointer.h"
3
4 MyIntPtr::MyIntPtr(int* value): value_{value} {}
5
6 int* MyIntPtr::GetValue() {
7     return value_
8 }
9
10 MyIntPtr::~~MyIntPtr() {
11     // Similar to C's free function.
12     delete value_;
13 }
```

```
1 void fn() {
2     // Similar to C's malloc
3     MyIntPtr p{new int{5}};
4     // Copy the pointer;
5     MyIntPtr q{p.GetValue()};
6     // p and q are both now destructed.
7     // What happens?
8 }
```

Smart Pointers

- Ways of wrapping unnamed (i.e. raw pointer) heap objects in named stack objects so that object lifetimes can be managed much easier
- Introduced in C++11
- Usually two ways of approaching problems:
 - `unique_ptr` + raw pointers ("observers")
 - `shared_ptr` + `weak_ptr`/raw pointers

Type	Shared ownership	Take ownership
<code>std::unique_ptr<T></code>	No	Yes
raw pointers	No	No
<code>std::shared_ptr<T></code>	Yes	Yes
<code>std::weak_ptr<T></code>	No	No

Unique pointer

- **std::unique_ptr<T>**
 - The unique pointer owns the object
 - When the unique pointer is destructed, the underlying object is too
- **raw pointer (observer)**
 - Unique Ptr may have many observers
 - This is an appropriate use of raw pointers (or references) in C++
 - Once the original pointer is destructed, you must ensure you don't access the raw pointers (no checks exist)
 - These observers **do not** have ownership of the pointer

Also note the use of 'nullptr' in C++ instead of NULL

Unique pointer: Usage

```
1 #include <memory>
2 #include <iostream>
3
4 int main() {
5     std::unique_ptr<int> up1{new int};
6     std::unique_ptr<int> up2 = up1; // no copy constructor
7     std::unique_ptr<int> up3;
8     up3 = up2; // no copy assignment
9
10    up3.reset(up1.release()); // OK
11    std::unique_ptr<int> up4 = std::move(up3); // OK
12    std::cout << up4.get() << "\n";
13    std::cout << *up4 << "\n";
14    std::cout << *up1 << "\n";
15 }
```

Can we remove "new" completely?

Observer Ptr: Usage

```
1 #include <memory>
2 #include <iostream>
3
4 int main() {
5     std::unique_ptr<int> up1(new int{0});
6     *up1 = 5;
7     std::cout << *up1 << "\n";
8     int* op1 = up1.get();
9     *op1 = 6;
10    std::cout << *op1 << "\n";
11    up1.reset();
12    std::cout << *op1 << "\n";
13 }
```

Unique Ptr Operators

```
1 #include <memory>
2 #include <iostream>
3
4 int main() {
5     // 1 - Worst - you can accidentally own the resource multiple
6     // times, or easily forget to own it.
7     int *i = new int;
8     auto up1 = std::make_unique<std::string>(i);
9     auto up11 = std::make_unique<std::string>(i);
10
11     // 2 - Not good - requires actual thinking about whether there's a leak.
12     std::unique_ptr<std::string> up2{new std::string{"Hello"}};
13
14     // 3 - Good - no thinking required.
15     std::unique_ptr<std::string> up3 = std::make_unique<std::string>("Hello");
16
17     std::cout << *up3 << "\n";
18     std::cout << *(up3.get()) << "\n";
19     std::cout << up3->size();
20 }
```

- <https://stackoverflow.com/questions/37514509/advantages-of-using-stdmake-unique-over-new-operator>
- <https://stackoverflow.com/questions/20895648/difference-in-make-shared-and-normal-shared-ptr-in-c>

Shared pointer

- **std::shared_ptr<T>**
- Several shared pointers share ownership of the object
 - A reference counted pointer
 - When a shared pointer is destructed, **if it is the only shared pointer left** pointing at the object, then the **object is destroyed**
 - May also have many observers
 - Just because the pointer has shared ownership doesn't mean the observers should get ownership too - don't mindlessly copy it
- **std::weak_ptr<T>**
 - Weak pointers are used with share pointers when:
 - You don't want to add to the reference count
 - You want to be able to check if the underlying data is still valid before using it.

Shared pointer: Usage

```
1 #include <memory>
2 #include <iostream>
3
4 int main() {
5     std::shared_ptr<int> x(new int{5});
6     std::shared_ptr<int> y = x; // Both now own the memory
7     std::cout << "use count: " << x.use_count() << "\n";
8     std::cout << "value: " << *x << "\n";
9     x.reset(); // Memory still exists, due to y.
10    std::cout << "use count: " << y.use_count() << "\n";
11    std::cout << "value: " << *y << "\n";
12    y.reset(); // Deletes the memory, since
13    // no one else owns the memory
14    std::cout << "use count: " << x.use_count() << "\n";
15    std::cout << "value: " << *y << "\n";
16 }
```

Can we remove "new" completely?

Weak Pointer: Usage

```
1 #include <memory>
2 #include <iostream>
3
4 int main() {
5     std::shared_ptr<int> x = std::make_shared<int>(1);
6     std::weak_ptr<int> wp = x; // x owns the memory
7     {
8         std::shared_ptr<int> y = wp.lock(); // x and y own the memory
9         if (y) {
10             // Do something with y
11             std::cout << "Attempt 1: " << *y << '\n';
12         }
13     } // y is destroyed. Memory is owned by x
14     x.reset(); // Memory is deleted
15     std::shared_ptr<int> z = wp.lock(); // Memory gone; get null ptr
16     if (z) {
17         // will not execute this
18         std::cout << "Attempt 2: " << *z << '\n';
19     }
20 }
```

When to use which type

- **Unique pointer vs shared pointer**
 - You almost always want a unique pointer over a shared pointer
 - Use a shared pointer if either:
 - An object has multiple owners, **and you don't know which one will stay around the longest**
 - You need temporary ownership (outside scope of this course)
 - This is very rare

When to use which type

- **Let's look at an example:**
 - `//lectures/week5/reader.cpp`

Shared or unique pointer?

- Computing examples
 - Linked list
 - Doubly linked list
 - Tree
 - DAG (mutable and non-mutable)
 - Graph (mutable and non-mutable)
 - Twitter feed with multiple sections (eg. my posts, popular posts)
- Real-world examples
 - The screen in this lecture theatre
 - The lights in this room
 - A hotel keycard
 - Lockers in a school

“Leak freedom in C++” poster

Strategy	Natural examples	Cost	Rough frequency
1. Prefer scoped lifetime by default (locals, members)	Local and member objects – directly owned	Zero: Tied directly to another lifetime	<div> <div>O(80%) of objects</div> <div>O(20%) of objects</div> </div>
2. Else prefer make_unique & unique_ptr or a container, if the object must have its own lifetime (i.e., heap) and ownership can be unique w/o owning cycles	Implementations of trees, lists	Same as new/delete & malloc/free Automates simple heap use in a library	
3. Else prefer make_shared & shared_ptr , if the object must have its own lifetime (i.e., heap) and shared ownership w/o owning cycles	Node-based DAGs, incl. trees that share out references	Same as manual reference counting (RC) Automates shared object use in a library	

Don't use owning raw *'s == don't use explicit *delete*

Don't create ownership cycles across modules by owning “upward” (violates layering)

Use *weak_ptr* to break cycles

Stack unwinding

- Stack unwinding is the process of exiting the stack frames until we find an exception handler for the function
- This calls any destructors on the way out
 - Any resources not managed by destructors won't get freed up
 - If an exception is thrown during stack unwinding, `std::terminate` is called

Not safe

```
1 void g() {  
2     throw std::runtime_error{""};  
3 }  
4  
5 int main() {  
6     auto ptr = new int{5};  
7     g();  
8     // Never executed.  
9     delete ptr;  
10 }
```

Safe

```
1 void g() {  
2     throw std::runtime_error{""};  
3 }  
4  
5 int main() {  
6     auto ptr = std::make_unique<int>(5);  
7     g();  
8 }
```

Exceptions & Destructors

- During stack unwinding, `std::terminate()` will be called if an exception leaves a destructor
- The resources may not be released properly if an exception leaves a destructor
- All exceptions that occur inside a destructor should be handled inside the destructor
- Destructors usually don't throw, and need to explicitly opt in to throwing
 - STL types don't do that

RAII

- Resource acquisition is initialisation
- A concept where we encapsulate resources inside objects
 - Acquire the resource in the constructor
 - Release the resource in the destructor
 - eg. Memory, locks, files
- Every resource should be owned by either:
 - Another resource (eg. smart pointer, data member)
 - The stack
 - A nameless temporary variable

Partial construction

Spot the bug

- What happens if an exception is thrown halfway through a constructor?
 - The C++ standard: "An object that is partially constructed or partially destroyed will have destructors executed for all of its fully constructed subobjects"
 - A destructor is not called for an object that was partially constructed
 - Except for an exception thrown in a constructor that delegates (why?)

```
1 #include <exception>
2
3 class my_int {
4 public:
5     my_int(int const i) : i_{i} {
6         if (i == 2) {
7             throw std::exception();
8         }
9     }
10 private:
11     int i_;
12 };
13
14 class unsafe_class {
15 public:
16     unsafe_class(int a, int b)
17         : a_{new my_int{a}}
18         , b_{new my_int{b}}
19     {}
20
21     ~unsafe_class() {
22         delete a_;
23         delete b_;
24     }
25 private:
26     my_int* a_;
27     my_int* b_;
28 };
29
30 int main() {
31     auto a = unsafe_class(1, 2);
32 }
```

Partial construction: Solution

- Option 1: Try / catch in the constructor
 - Very messy, but works (if you get it right...)
 - Doesn't work with initialiser lists (needs to be in the body)
- Option 2:
 - An object managing a resource should initialise the resource last
 - The resource is only initialised when the whole object is
 - Consequence: An object can only manage one resource
 - If you want to manage multiple resources, instead manage several wrappers , which each manage one resource

```
1 #include <exception>
2 #include <memory>
3
4 class my_int {
5 public:
6     my_int(int const i)
7         : i_{i} {
8         if (i == 2) {
9             throw std::exception();
10        }
11    }
12 private:
13     int i_;
14 };
15
16 class safe_class {
17 public:
18     safe_class(int a, int b)
19         : a_(std::make_unique<my_int>(a))
20         , b_(std::make_unique<my_int>(b))
21     {}
22 private:
23     std::unique_ptr<my_int> a_;
24     std::unique_ptr<my_int> b_;
25 };
26
27 int main() {
28     auto a = safe_class(1, 2);
29 }
```

Reference types

Reference types

We learnt about *references* to *objects* in Week 1.

Reference types

We learnt about *references* to *objects* in Week 1.

We learnt about *containers*, *iterators*, and *views* in Week 2.

Reference types

We learnt about *references* to *objects* in Week 1.

We learnt about *containers*, *iterators*, and *views* in Week 2.

We've just learnt about *smart pointers* in Week 4.

Reference types

We learnt about *references* to *objects* in Week 1.

We learnt about *containers*, *iterators*, and *views* in Week 2.

We've just learnt about *smart pointers* in Week 4.

How do they all tie together and what's the relationship with *pointers*?

Reference types

A *reference type* is a type that acts as an abstraction over a raw pointer.

They don't own a resource, but allow us to do cheap operations like copy, move, and destroy (see Week 5) on objects that do.

A reference type to a range of elements is called a *view*.

```
class string_view6771 {
public:
    string_view6771() noexcept = default;

    explicit(false) string_view6771(std::string const& s) noexcept
    : string_view6771(s.data())
    {}

    explicit(false) string_view6771(char const* data) noexcept
    : string_view6771(data, std::strlen(data))
    {}

    string_view6771(char const* data, std::size_t const length) noexcept
    : data_{data}
    , length_{length}
    {}

    auto begin() const noexcept -> char const* { return data_; }
    auto end() const noexcept -> char const* { return data_ + length_; }
    auto size() const noexcept -> std::size_t { return length_; }
    auto data() const noexcept -> char const* { return data_; }

    auto operator[](std::size_t const n) const noexcept -> char {
        assert(n < length_);
        return data_[n];
    }
private:
    char const* data_ = nullptr;
    std::size_t length_ = 0;
};
```

Reference types

`T*`

`T const*`

Raw pointer to a single object or to an element in a range. Wherever possible, prefer references, iterators, or something below.

`std::string_view`

Abstraction over **immutable** string-like data.

`std::span<T>`

`std::span<T const>`

Abstraction over array-like data.

Most

`ranges::views::*`

Lazy abstraction over a range, associated with some transformation.

A quick look at span

```
1 auto zero_out(std::span<int> const x) -> void {
2     ranges::fill(x, 0);
3 }
4
5 {
6     auto numbers = views::iota(0, 100) | ranges::to<std::vector>;
7     zero_out(numbers);
8     CHECK(ranges::all_of(numbers, [](int const x) { return x == 0; }));
9 }
10 {
11     // Using int[] since spec requires we use T[] instead of std::vector
12     // NOLINTNEXTLINE(modernize-avoid-c-arrays)
13     auto const raw_data = std::make_unique<int[]>(42);
14     auto const usable_data = std::span<int>(raw_data.get(), 42);
15
16     zero_out(usable_data);
17     CHECK(ranges::all_of(usable_data, [](int const x) { return x == 0; }));
18 }
```

A quick look at span

```
1 auto zero_out(std::span<int> const x) -> void {
2     ranges::fill(x, 0);
3 }
4
5 {
6     auto numbers = views::iota(0, 100) | ranges::to<std::vector>;
7     zero_out(numbers);
8     CHECK(ranges::all_of(numbers, [](int const x) { return x == 0; }));
9 }
10 {
11     // Using int[] since spec requires we use T[] instead of std::vector
12     // NOLINTNEXTLINE(modernize-avoid-c-arrays)
13     auto const raw_data = std::make_unique<int[]>(42);
14     auto const usable_data = std::span<int>(raw_data.get(), 42);
15
16     zero_out(usable_data);
17     CHECK(ranges::all_of(usable_data, [](int const x) { return x == 0; }));
18 }
```

NOLINTNEXTLINE(check) turns off a check on the following line.

A quick look at span

```
1 auto zero_out(std::span<int> const x) -> void {
2     ranges::fill(x, 0);
3 }
4
5 {
6     auto numbers = views::iota(0, 100) | ranges::to<std::vector>;
7     zero_out(numbers);
8     CHECK(ranges::all_of(numbers, [](int const x) { return x == 0; }));
9 }
10 {
11     // Using int[] since spec requires we use T[] instead of std::vector
12     // NOLINTNEXTLINE(modernize-avoid-c-arrays)
13     auto const raw_data = std::make_unique<int[]>(42);
14     auto const usable_data = std::span<int>(raw_data.get(), 42);
15
16     zero_out(usable_data);
17     CHECK(ranges::all_of(usable_data, [](int const x) { return x == 0; }));
18 }
```

NOLINTNEXTLINE(check) turns off a check on the following line.

Do this rarely, and always provide justification immediately above.

A quick look at span

```
1 auto zero_exists(std::span<int const> const x) -> bool {
2     ranges::any_of(x, [](int const x) { return x == 0; });
3 }
4
5 {
6     auto numbers = views::iota(0, 100) | ranges::to<std::vector>;
7     CHECK(zero_exists(numbers));
8 }
9 {
10    // Using int[] since ass2 spec requires we use T[] instead of std::vector
11    // NOLINTNEXTLINE(modernize-avoid-c-arrays)
12    auto const raw_data = std::make_unique<int[]>(42);
13    auto const usable_data = std::span<int>(raw_data.get(), 42);
14
15    CHECK(zero_exists(usable_data));
16 }
```

A quick look at span

```
1 auto zero_exists(std::span<int const> const x) -> bool {
2     ranges::any_of(x, [](int const x) { return x == 0; });
3 }
4
5 {
6     auto numbers = views::iota(0, 100) | ranges::to<std::vector>;
7     CHECK(zero_exists(numbers));
8 }
9 {
10    // Using int[] since ass2 spec requires we use T[] instead of std::vector
11    // NOLINTNEXTLINE(modernize-avoid-c-arrays)
12    auto const raw_data = std::make_unique<int[]>(42);
13    auto const usable_data = std::span<int>(raw_data.get(), 42);
14
15    CHECK(zero_exists(usable_data));
16 }
```

NOLINTNEXTLINE(check) turns off a check on the following line.

A quick look at span

```
1 auto zero_exists(std::span<int const> const x) -> bool {
2     ranges::any_of(x, [](int const x) { return x == 0; });
3 }
4
5 {
6     auto numbers = views::iota(0, 100) | ranges::to<std::vector>;
7     CHECK(zero_exists(numbers));
8 }
9 {
10    // Using int[] since ass2 spec requires we use T[] instead of std::vector
11    // NOLINTNEXTLINE(modernize-avoid-c-arrays)
12    auto const raw_data = std::make_unique<int[]>(42);
13    auto const usable_data = std::span<int>(raw_data.get(), 42);
14
15    CHECK(zero_exists(usable_data));
16 }
```

NOLINTNEXTLINE(check) turns off a check on the following line.

Do this rarely, and always provide justification immediately above.

Iterators and pointers

“ A pointer is an abstraction of a virtual memory address.

—Sy Brand

Iterators and pointers

“ A pointer is an abstraction of a virtual memory address.

—Sy Brand

“ Iterators are a family of concepts that abstract different aspects of addresses, ...

—Elements of Programming, Stepanov & McJones

Iterators and pointers

“ A pointer is an abstraction of a virtual memory address.

—Sy Brand

“ Iterators are a family of concepts that abstract different aspects of addresses, ...

—Elements of Programming, Stepanov & McJones

“ Iterators are a generalization of pointers that allow a C++ program to work with different data structures... in a uniform manner.

—Working Draft, Standard for Programming Language C++

Ownership and lifetime with reference types

We need to be very careful when using reference types.

The owner must always outlive the observer.

```
auto v = std::vector<int>{0, 1, 2, 3};
auto const* p = v.data();

{
    CHECK(*p == 0); // okay: p points to memory
                    // owned by v
}

v = std::vector<int>{0, 1, 2, 3};

{
    CHECK(*p == 0); // error: p points to memory
                    // owned by no one
}
```

Ownership and lifetime with reference types

We need to be very careful when using reference types.

The owner must always outlive the observer.

```
auto not_okay1(std::string s) -> std::string_view { // s is a local; will be destroyed
    return s;                                     // before its observer
} // s destroyed here

auto not_okay2(std::string const& s) -> std::string_view { // s may be destroyed before observer;
    return s;                                             // considered harmful
}                                                         // e.g. auto x = not_okay("hello")

auto okay1(std::string_view const sv) -> std::string_view { // observer in, observer out
    return sv;
}

auto okay2(std::string& s) -> std::string_view { // lvalue reference in, observer out
    return s;
}
```


Ownership and lifetime with reference types

We need to be very careful when using reference types.

The owner must always outlive the observer.

```
auto is_even(int const x) -> bool {
    return x % 2 == 0;
};

auto not_okay1(std::vector<int> v) {
    return v | views::filter(is_even);
} // v destroyed here

auto not_okay2(std::vector<int> const& v) {
    return v | views::filter(is_even);
}
```

```
auto okay1(std::span<int> s) {
    return s | views::filter(is_even);
}

auto okay2(std::vector<int>& v) {
    return v | views::filter(is_even);
}
```